Structural Evolution and User Migration in Dynamic Interest Graphs: A Case Study of MyAnimeList (2006–2018)

Iaroslav Sagan (66661)
University of Lisbon
Lisbon, Portugal
yaroslav_sagan@gmail.com

Anna Maksymchuk (66662)
University of Lisbon
Lisbon, Portugal
anna.i.maksymchuk@gmail.com

Maria Samosudova (66663)
University of Lisbon
Lisbon, Portugal
samosudova@gmail.com

ABSTRACT

This study analyzes the structural evolution of the MyAnimeList platform (2006–2018). Utilizing time-sliced networks, we examine the dynamics of taste communities and user navigation. We employ the Leiden algorithm to trace the genealogy of clusters from a cohesive core to a fragmented landscape. Regarding user modeling, we find that stochastic random walker agents diverge from empirical trajectories, indicating the limitations of purely topological simulation. In contrast, machine learning experiments successfully predict user migration patterns. Crucially, we demonstrate that the inclusion of community-based clustering metrics significantly improves model accuracy, highlighting the predictive value of mesoscale network features.

1 INTRODUCTION

1.1 The MyAnimeList Ecosystem as a Complex Network

The digital aggregation of cultural preferences has transformed the study of computational sociology. MyAnimeList (MAL), established in 2006, serves as a primary registry for anime consumption. Unlike generic social networks where edges represent declared friendships, the primary structure of MAL is a bipartite interest graph connecting users to titles. This structure provides an optimal environment for analyzing "affinity networks" [3], where community formation is driven by taste homophily rather than geographic proximity.

The period between 2006 and 2018 represents a critical epoch in the globalization of Japanese animation, characterized by the transition from a niche subculture to a mainstream entertainment force. Analyzing this period allows for the observation of a complete evolutionary cycle of a complex network—from a cohesive core to a fragmented, multi-polar topology.

1.2 Objectives and Research Questions

The primary objective of this study is to model the MAL community as a dynamic, evolving system. Traditional static graph analysis fails to capture the temporal fluidity of user interests. To address this, we employ a multi-stage analytical framework combining dynamic community detection, stochastic simulation, and predictive modeling.

A key focus of our investigation is the evaluation of topological agents. We test the hypothesis that standard Random Walker models [6], which rely solely on structural connectivity, are sufficient to simulate user navigation in affinity networks. Furthermore, we explore whether the predictive power of machine learning models can be enhanced by incorporating mesoscale network features derived from community detection.

The research is guided by the following key questions:

- **Topological Evolution:** How did the structural properties of the MAL graph change during the community's expansion (2006-2018)?
- **Community Dynamics:** How do taste clusters evolve over time, and can we identify distinct phases of fragmentation using dynamic community detection?
- **Limits of Topological Simulation:** To what extent do empirical user trajectories diverge from theoretical Random Walk models?
- **Predictive Modeling:** Can we predict user migration between communities using supervised learning, and does the inclusion of clustering metrics improve model performance?

2 DATA AND METHODOLOGY

2.1 Dataset Acquisition and Preprocessing

The data originates from a publicly available kaggle datasets aggregated from MAL profiles [1, 11], covering the period from 2006 to 2018. Given the platform's predominantly international user base, the dataset reflects global consumption patterns rather than domestic Japanese trends.

To ensure data quality, we applied a multi-stage filtering pipeline:

- **Bot Detection:** Removal of inactive accounts and profiles exhibiting automated behavior.
- **Percentile-based Truncation:** We excluded users falling into the extreme tails of the activity distribution. This removes users with too few votes (insufficient signal for clustering) and those with implausibly high vote counts, ensuring the analysis focuses on human-scale consumption patterns.

The final processed dataset comprises approximately **85,000 unique users** and **6,500 anime titles**.

2.2 Graph Projection and Topology Construction

We model the system as a bipartite graph which is subsequently projected into two distinct monopartite networks. The detailed construction pipelines are described in Sections **??** and **??**, respectively.

2.2.1 Anime-Anime Network. In this projection, an edge exists if two titles share a common voter. To account for varying audience sizes, we utilized the **Jaccard Similarity** [8] index as the edge weight.

$$J(A, B) = \frac{|U_A \cap U_B|}{|U_A \cup U_B|}$$

Given the extreme density of the raw projection (where a single popular anime could fully connect thousands of users), we applied

a hard threshold of $J > 0.05$. This effectively prunes weak links formed by random coincidences while preserving significant genre or fandom connections.

2.2.2 User-User Network. Here, an edge connects two users if they have rated the same anime. The edge weight is defined as the raw count of shared titles (co-votes). A major challenge in this projection is the variance in edge weights, which range from negligible values (2-3 shared items) to tens of thousands (10^4). To address this, we implemented a cutoff threshold: edges were retained only if users shared more than **3 titles**.

Anyway, even after thresholding, the user-user raw network projection suffered from extreme density saturation. Popular "blockbuster" titles (e.g., *Death Note*, *Attack On Titan*) act as super-hubs; a single vote for such a title effectively connects a user to thousands of others, creating a near-clique structure that obscures genuine taste communities. So, this titles had to be deleted from the dataset prior to projection.

2.2.3 Further Sparsification Attempts. More aggressive sparsification techniques (e.g., Backbone extraction, k-NN) were tested but did not reveal significantly distinct structural patterns. Consequently, we retained the simpler approach to avoid unnecessary information loss while maintaining structural clarity.

2.3 Resulting Topology

These thresholding strategies proved effective in mitigating the "hairball" phenomenon common in social graphs. The resulting networks exhibited a graph density in the range of $0.2 - 0.3$, striking a balance between sparsity (for efficient clustering) and connectivity (preserving the Giant Connected Component).

Only after these topological corrections is the graph subjected to the Leiden community detection algorithm and Random Walk simulations.

2.4 Computational Framework and Reproducibility

To ensure reproducibility and handle large-scale temporal networks efficiently, we developed a dedicated modular Python framework `project_cda` [7], which is a core part of the open-source repository MARS_1.0. The framework is designed with a modular architecture to support the full research lifecycle:

Graph Construction Engine. The graph construction modules `AnimeGraphBuilder` and `UserGraphBuilder` modules utilize streaming JSON parsers (‘ijson‘) to process massive user interaction logs with minimal memory footprint. They implement the projection logic described in Section 3.1 and support vectorized graph operations via the ‘igraph‘ C-core [4], ensuring high performance even for dense snapshots.

Simulation and Analysis Modules. The framework includes specialized components for dynamic analysis:

- **CommunityTracker:** Implements a greedy matching algorithm based on Jaccard similarity to trace cluster lineage across time steps.
- **RandomCrowd:** An agent-based simulation engine that deploys stochastic walkers to probe network topology and user navigation patterns.
- **ClusterEvaluation:** Encapsulates the calculation of structural (Modularity), semantic (Purity), and information-theoretic (Entropy) metrics.

Data Management and Caching. Given the computational cost of generating 13 annual snapshots with varying hyperparameters, we implemented a robust caching system managed by a `PathManager`. This module enforces idempotency: each experimental configuration (edge weights, sparsification, clustering algorithm) generates a unique hash signature. If a serialized artifact exists for a given configuration, it is loaded instantly ("lazy evaluation"), preventing redundant computations. The exact data schema required to run the pipeline is documented in the repository's `README.md`.

3 TOPOLOGICAL EVOLUTION OF PROJECTED NETWORKS

In this section[1], we analyze the structural evolution of the projections derived from the bipartite graph. We first examine the Anime-Anime network to understand how content relationships shifted over time, followed by an analysis of the User-User network.

3.1 Anime-Anime Network

The projected anime-anime network experienced explosive growth over the analyzed period (2006–2018), transforming from a compact, niche community into a sprawling, heterogeneous ecosystem. This transformation is defined by three primary phenomena: the densification-sparsification paradox, increasing taste divergence, and the crystallization of a "rich-club" core.

3.1.1 The Densification-Sparsification Paradox. The network underwent a dramatic scale expansion: the number of nodes (anime titles) increased from 732 in 2006 to 6,129 in 2018, while the volume of connections (edges) surged from ∼64,000 to ∼819,000. However, this volumetric growth reveals a fundamental structural shift.

While the absolute number of connections increased by an order of magnitude, the potential number of connections grew quadratically (N^2). Consequently, the global Graph Density declined precipitously from 0.2387 (2006) to 0.0436 (2018).

This indicates a transition from a "Village" topology—where the community is small enough for high interconnectedness—to a "Metropolis" structure. In the modern era, the ecosystem has become highly specialized; while the total volume of interactions is higher, individual anime titles connect to a significantly smaller fraction of the total population. The network has shifted from a monolithic block to a spread-out, sparse landscape.

3.1.2 Increasing Social Distance and Taste Divergence. To quantify the "cost" of traversing this expanding network, we analyzed weighted path metrics. Since edge weights represent similarity (Jaccard), the weighted distance can be interpreted as "social distance" or taste divergence.

[1]The complete experimental pipeline and visualization tools are available in the project repository: `project_cda/1_general_graph_metrics.ipynb`.

The evolution of these metrics is presented in **Figure 1**. As shown in the *upper-left panel*, the average weighted path length rose sharply from 7.1 in 2006 to 44.4 in 2018. This metric represents the "resistance" to navigation: connecting a fan of a niche genre to a mainstream hit now requires passing through significantly more intermediaries.

Simultaneously, the network diameter (*upper-right panel*) expanded from 29 to 188.5 weighted units. This confirms that the "taste universe" is expanding. Distinct clusters (e.g., modern idols vs. vintage mecha) are moving mathematically further apart, creating deep topological fissures.

3.1.3 Local Cohesion and the "Fandom" Effect. Despite the global sparsification, the network maintains robust local connectivity. The *bottom-left panel* of **Figure 1** illustrates the Average Clustering Coefficient. After an initial adjustment, the metric stabilized at a remarkably high value of ≈0.59. This indicates that the "Small-World" property [9] is preserved locally. If Anime A is connected to B and C, there is a consistent ∼60% probability that B and C are also connected. This proves that the sparsification did not destroy community cohesion; instead, the landscape fractured into tight, self-reinforcing "genre bubbles" (fandoms).

3.1.4 Structural Phase Shift: The 2006 Anomaly. The year 2006 represents a distinct topological phase. In this nascent period, the network exhibited disassortative mixing (Degree Assortativity ≈ −0.11), suggesting a star-like structure where popular titles served as hubs connecting primarily to niche nodes. From 2007 onward, the network flipped to positive assortativity (≈ 0.50), signaling the emergence of the "Rich-Club" phenomenon.

3.1.5 Intensity vs. Topology. Finally, we examine the distribution of influence using Node Strength. The log-log plot in the *bottom-right panel* of **Figure 1** confirms the scale-free nature of the modern network, following a clear Power Law distribution. The network is dominated by a few "mega-hubs," validating the "preferential attachment" growth model.

However, the Assortativity of Strength (≈ 0.18) is consistently lower than the Assortativity of Degree (≈ 0.50). This implies that while popular shows are structurally connected, the strongest taste affinities are located in the niche clusters, not the mainstream core.

3.2 User-User Network

In stark contrast to the Anime content network—which became "sparse" and harder to traverse as it grew—the User interaction network exhibits the classic properties of Network Densification. As the community expanded, the social distance between users collapsed, making the network significantly more interconnected.

While the "universe" of users grew, the social structure did not fragment into isolated islands. Instead, it evolved into a tight, integrated "global village," where new users actively connected to existing hubs rather than the periphery.

3.2.1 Global Integration and the "Shrinking World". The analysis of weighted path metrics reveals a community that is becoming functionally smaller and easier to traverse, despite growing in physical size.

The evolution of these metrics is presented in the *upper panels* of **Figure 2**. As shown in the *upper-left panel*, the average weighted path length dropped sharply from ∼0.45 in 2006 to ∼0.31 by 2009, maintaining this lower baseline through 2018. This reduction is a hallmark of the "Small World" effect: as the platform matured, users formed bridging connections, accelerating the flow of information across the graph.

Similarly, the network diameter (*upper-right panel*) contracted from ∼1.08 to < 0.99. Unlike the Anime graph, where taste divergence created massive gaps, the social graph's diameter is shrinking. This indicates that even socially distant groups (e.g., distinct language communities) are becoming more connected to the mainstream core.

3.2.2 Local Cohesion and Community Structure. While the network became globally smaller, the local structure evolved to balance rapid growth with intimate social circles. The *bottom-left panel* of **Figure 2** shows the Average Clustering Coefficient. It peaked in 2009 (∼0.73) during the platform's initial boom, followed by a gentle decline to ∼0.66.

A score of 0.66 remains exceptionally high for a large social network. The slight decline suggests a natural dilution of cliques as users added diverse friends, but the high retention proves the community is fundamentally built on strong, overlapping friend groups rather than loose acquaintances.

3.2.3 Influence and Inequality. The distribution of influence confirms a highly stratified social hierarchy. The log-log plot in the *bottom-right panel* demonstrates a strict linear descent, characteristic of a Scale-Free Network ($P(k) \sim k^{-\gamma}$) [2].

The graph is dominated by a tiny fraction of "Super-Users" (hubs) who possess nearly $1,000\times$ the connectivity of the average user. These hubs act as the structural "glue" that holds the giant component together, enabling the short path lengths observed in Section 3.2.

4 MESOSCALE ANALYSIS AND COMMUNITY DETECTION EXPERIMENTS

Having characterized[2] the global topological evolution, we shift our focus to the mesoscale level—identifying the distinct taste communities that constitute the anime ecosystem. This section details the experimental framework used to select the optimal clustering algorithm, ensuring that the detected communities are structurally robust, semantically meaningful, and suitable for the subsequent migration analysis.

4.1 Experimental Setup

For each annual snapshot G_t of the projected Anime-Anime network (2006–2018), we conducted a comparative analysis of five primary community detection algorithms:

(1) **Leiden (Modularity):** Tested with resolution parameters $\gamma \in [0.5, 2.0]$.
(2) **Infomap (Map Equation):** Tested with Markov time parameters $T \in \{10, 25, 50\}$.

[2]The complete experimental pipeline and visualization tools are available in the project repository: `project_cda/2_clusterization.ipynb`.

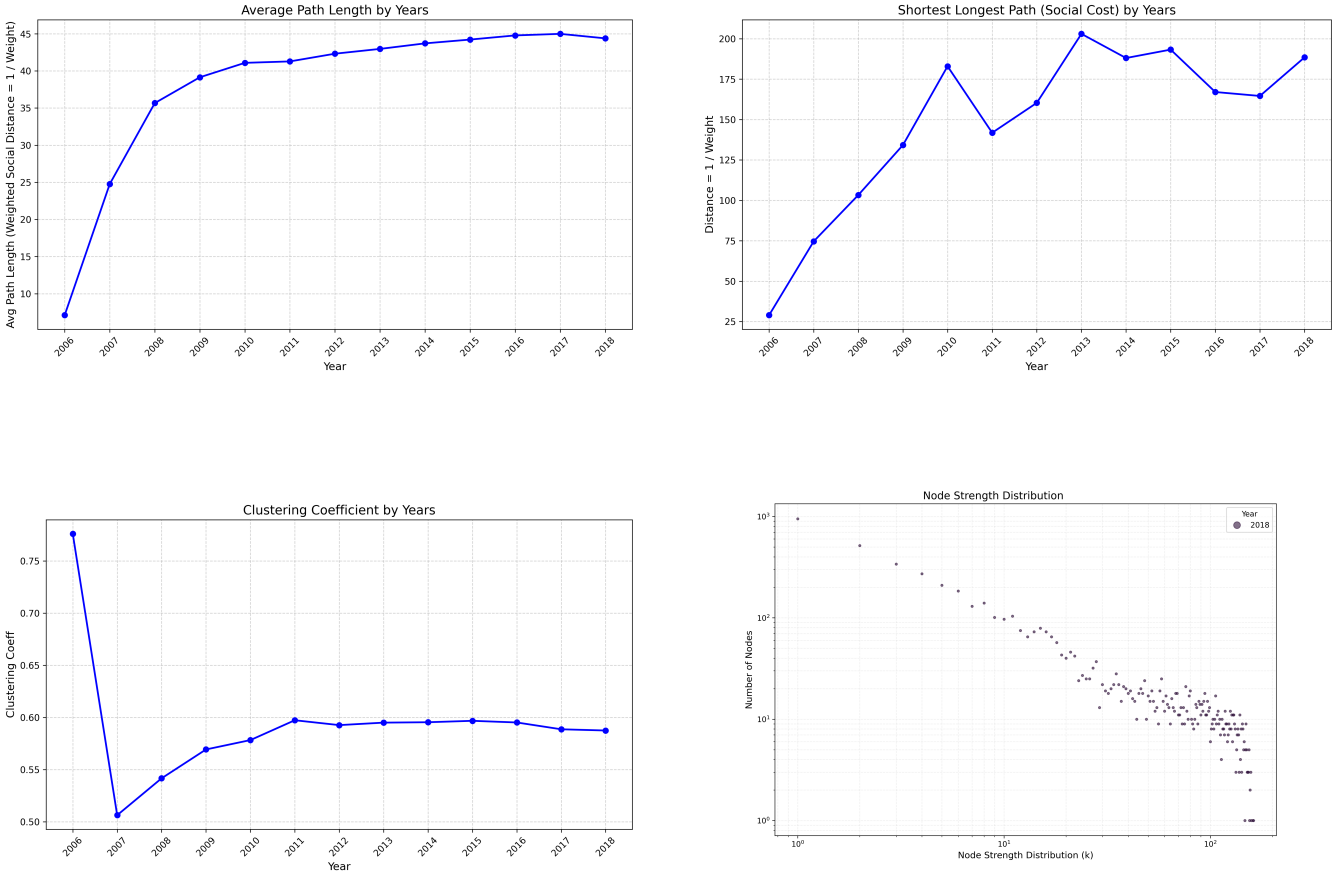

Figure 1: Evolution of Anime Network Topology (2006–2018). Upper-left: Average Weighted Path Length showing increased navigation difficulty. Upper-right: Network Diameter indicating the expansion of the "taste universe". Bottom-left: Average Clustering Coefficient stabilizing around 0.59, suggesting persistent local cohesion. Bottom-right: Node Strength Distribution (2018) confirming the scale-free ($P(k) \sim k^{-\gamma}$) nature of the modern network.

(3) **Leading Eigenvector:** A spectral method based on modularity maximization.
(4) **Label Propagation (LPA):** Used as a baseline heuristic.

The quality of partitions was evaluated using a composite set of metrics, prioritizing **Modularity (Q)** for structural definition, **Genre Purity** for semantic alignment, and **Number of Clusters (N_{cl})** for interpretability.

4.2 Comparative Analysis: The Structure-Granularity Trade-off

The experiments revealed a clear trade-off between semantic precision and structural compactness (see Table 1).

Label Propagation (LPA) proved to be an outlier with suboptimal performance. It demonstrated the lowest Modularity ($Q \approx 0.03$), failing to detect the dense community structure known to exist in the network.

Leading Eigenvector showed extreme instability. While its average modularity (≈ 0.30) was acceptable, the number of clusters fluctuated wildly across years (ranging from 4 to 108), indicating high sensitivity to minor topological changes. Such volatility makes it unsuitable for longitudinal tracking.

Infomap exhibited behavior highly dependent on the Markov time parameter T:

- At $T = 10$ ("short walks"), it resulted in **hyper-fragmentation**, producing over 200 micro-communities with high semantic purity (~ 0.63) but poor modular structure ($Q \approx 0.19$).
- At $T = 50$, it achieved the highest semantic accuracy overall (Genre Purity ~ 0.61) while maintaining decent modularity ($Q \approx 0.35$). However, it still divided the network into ≈ 29 clusters, which is too granular for macro-level migration analysis.

Leiden (Modularity) demonstrated superior structural definition. It consistently achieved the highest Modularity scores ($Q \approx 0.36$). Crucially, at $\gamma = 1.0$, it struck a "Golden Mean":

- **High Modularity:** 0.358 (comparable to the best results).
- **Acceptable Purity:** 0.57 (close to Infomap's 0.61).

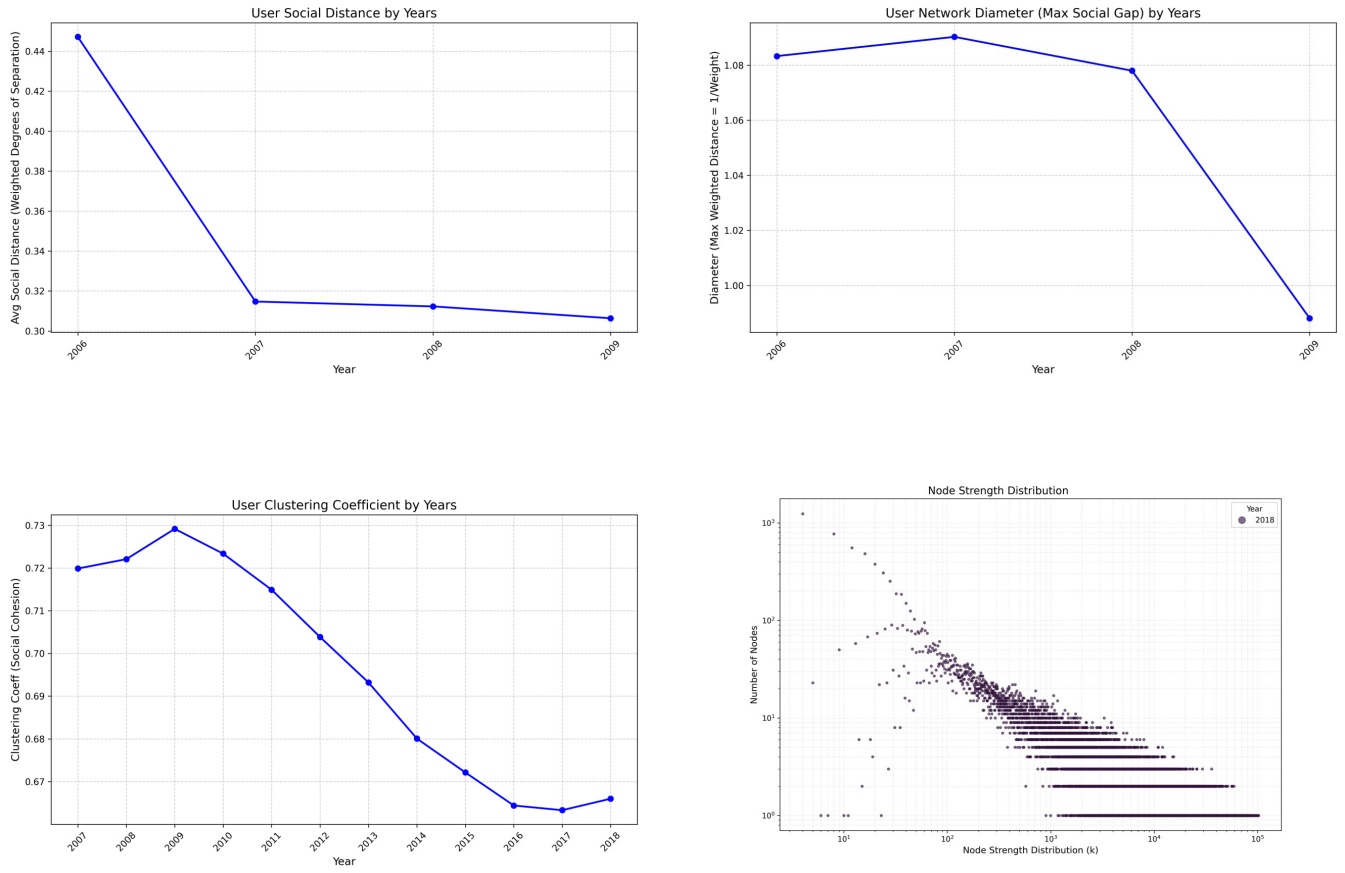

Figure 2: Evolution of User Network Topology (2006–2018). Upper-left: Avg. Weighted Path Length decreasing, indicating higher integration. Upper-right: Network Diameter contracting, showing the "shrinking world" phenomenon. Bottom-left: Clustering Coefficient peaking in 2009 and slowly stabilizing, reflecting the balance between clique formation and expansion. Bottom-right: Node Strength Distribution (2018) following a strict Power Law, highlighting the dominance of "Super-Users".

- **Optimal Compactness:** It identified $\approx$ 10 stable macro-communities.

4.3 Selection Logic

Based on these results, we prioritized structural interpretability over maximal semantic purity. For the task of predicting user migration, it is more valuable to track transitions between a manageable number of distinct "Macro-Communities" (e.g., "Mainstream Action" vs. "Vintage Mecha") rather than tracking noise across 100+ fragmented micro-clusters (as seen in Infomap T=10 or Eigenvector). Therefore, we selected **Leiden ($\gamma = 1.0$)** as the primary algorithm.

4.4 Resolution Parameter Sweep

Having selected the Leiden algorithm, we performed a detailed sensitivity analysis to confirm the stability of the chosen resolution $\gamma = 1.0$. We constructed heatmaps for Modularity and Cluster Counts across all years for $\gamma \in [0.1, 3.0]$.

As illustrated in **Figure 3**, the combined analysis confirms our choice:

- **Modularity Landscape:** We observe high modularity values persisting across the range $\gamma \in [0.8, 1.2]$, suggesting that the strong community structure is not an artifact of a specific parameter point.
- **Cluster Count Plateau:** Simultaneously, the number of clusters stabilizes in this range. Unlike higher resolutions ($\gamma > 1.5$), where the network splinters, $\gamma = 1.0$ consistently yields ~ 10 distinct communities.

This topological stability provides a solid foundation for the longitudinal tracking of user migration.

Table 1: Average Performance of Clustering Algorithms (2006–2018)

Algorithm	Modularity	Genre Purity	Source Purity	N Clusters	Verdict
Label Propagation	0.033	0.31	0.35	6	Underfitting
Leading Eigenvector	0.301	0.50	0.45	34	**Unstable**
Infomap ($T = 10$)	0.192	**0.63**	**0.54**	~ 150	Hyper-segmentation
Infomap ($T = 50$)	0.345	0.61	**0.54**	29	Over-segmentation
Leiden ($\gamma = 0.9$)	0.357	0.00*	0.00*	6	Semantic Mismatch
Leiden ($\gamma = 1.0$)	**0.358**	0.57	0.53	**10**	**Selected**

* Values of 0.00 indicate data corruption for specific years due to tag misalignment during evaluation.

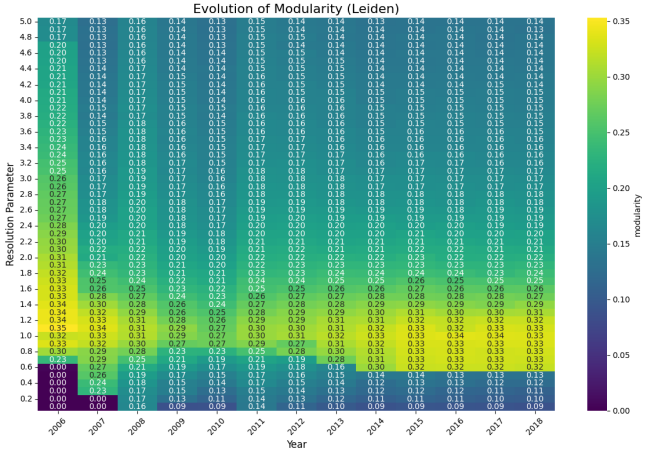

(a) Evolution of Modularity (Q).

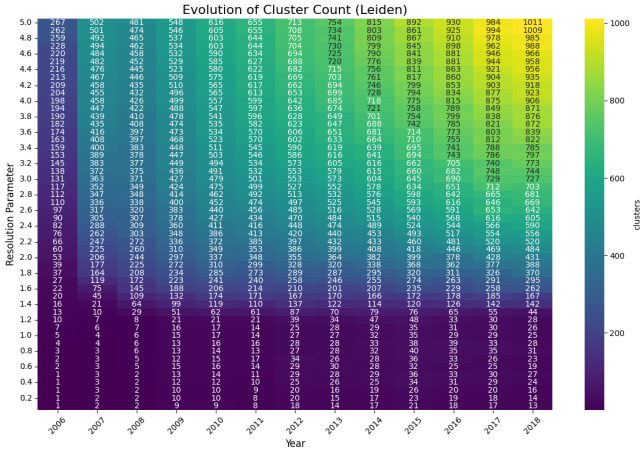

(b) Evolution of Cluster Count (N).

Figure 3: Multi-resolution analysis of the Leiden algorithm (2006–2018). (a) Modularity remains robust across the optimal range. (b) The number of clusters exhibits a stable "plateau" around $\gamma = 1.0$ (~ 10 communities).

5 STOCHASTIC SIMULATION OF USER NAVIGATION

In this chapter[3], we investigate whether user migration patterns can be explained by purely topological factors. To this end, we introduce a stochastic agent—the **Random Walker** — which serves as a null model for user navigation. By simulating thousands of probabilistic trajectories on the evolving graph, we test the "Structural Determinism" hypothesis: that a user's future community is primarily determined by the connectivity of their current position.

5.1 Formal Definition of the Topological Agent

Let $G_t = (V_t, E_t)$ denote the state of the network at year t. We model the user's journey as a sequence of community memberships:

$$\mathcal{T}_u = (c_0, c_1, \ldots, c_T)$$

where c_t is the community ID of user u at time t.

The Random Walker agent operates under the assumption of Markovian neutrality. Conditioned on being in node v at time t, the probability of transitioning to node w at $t + 1$ is proportional to the edge weight w_{vw}:

$$P(X_{t+1} = w \mid X_t = v) = \frac{w_{vw}}{\sum_{k \in \mathcal{N}(v)} w_{vk}} \tag{1}$$

This formulation implies that the agent has no memory and no specific preferences other than the structural "path of least resistance" offered by the graph topology.

5.2 Simulation Protocol

To generate a statistical baseline, we execute the following protocol for every user in the dataset:

(1) **Initialization:** An ensemble of $K = 100$ stochastic agents is spawned at the user's actual position in year $t = 2006$.
(2) **Propagation:** For each subsequent year $t + 1$, agents transition to a new community based on the transition matrix of the snapshot G_t.
(3) **Consensus:** The resulting cloud of K trajectories forms a probabilistic distribution of possible futures. We define the *Predicted Community* $\hat{c}_{t+1}$ as the mode of this distribution (Majority Vote).

This approach effectively turns the Random Walker into a classifier: if the user's actual next step coincides with the walker's most probable step, we consider the behavior "structurally predictable."

[3]The complete experimental pipeline and visualization tools are available in the project repository: `project_cda/3_random_walker.ipynb`.

5.3 Divergence from Empirical Trajectories

Comparing the simulated trajectories with empirical user data reveals a fundamental disconnect between topological probability and human choice. We quantified the similarity between the Random Walker's predicted path and the user's actual migration history using a population-wide similarity metric (Jaccard overlap of visited communities).

The results indicate a near-total divergence:

- **Mean Similarity:** 0.0435 (±0.0184)
- **Median Similarity:** 0.0447

These extremely low values (≈ 4.3%) imply that user navigation in the MAL ecosystem is **not stochastic**. While the network structure defines the **possibilities** (the set of accessible neighbors), it does not dictate the *probabilities*. The "gravity" of large structural hubs, which dominates the Random Walker's decisions, fails to explain the nuanced, content-driven choices made by real users.

5.4 Conclusion: The Need for Supervised Learning

The failure of the Random Walker experiment demonstrates that the MAL ecosystem is not merely a system of passive diffusion. Users are not particles flowing down topological gradients; their migration is driven by active choices, likely influenced by specific content features, social signals, and external trends not captured by edge weights alone.

This finding provides the motivation for the next chapter. Since simple Markovian dynamics are insufficient, we must turn to **Supervised Machine Learning** (Section 6) to capture the non-linear interactions between structural metrics and user behavior.

6 PREDICTIVE MODELING OF USER MIGRATION

In this chapter[4], we address the problem of user migration between communities, a fundamental challenge in dynamic social network analysis. We model the MAL user network as an evolving graph $G_t = (V_t, E_t)$, where nodes represent unique users and weighted edges denote social proximity based on shared anime consumption history. Since the statistical and topological properties of the network fluctuate significantly over time, our analysis focuses on transitions between consecutive temporal snapshots, denoted as t and $t + 1$.

The migration prediction task is formalized as a binary classification problem. Let $C(u, t)$ denote the community assignment of user u at time t. A user is labeled as a *migrant* if their community membership changes between snapshots:

$$y_u = \begin{cases} 1 & \text{if } C(u, t+1) \neq C(u, t) \\ 0 & \text{if } C(u, t+1) = C(u, t) \end{cases} \tag{2}$$

Features extracted at time t are utilized to predict the state y_u at time $t + 1$.

[4]The complete experimental pipeline and visualization tools are available in the project repository: `project_cda/4_user_migration_pipeline.ipynb`.

6.1 Motivation for Community Detection Algorithm Choice

To identify distinct user communities within each temporal snapshot, we employ the **Leiden algorithm**. While sharing the objective of modularity maximization with the classical Louvain method, Leiden is selected for its superior stability, scalability, and topological guarantees on large-scale networks.

The primary advantage of Leiden over Louvain lies in its handling of community connectivity [10], scalability and reliability on large networks. While Louvain is efficient, it can produce poorly connected or even disconnected communities, Leiden addresses this through a distinct *refinement phase*, guaranteeing that all resulting clusters are well-connected. This is critical for our study, as membership in a poorly connected cluster could represent an algorithmic artifact rather than genuine social cohesion.

Alternative community detection methods were evaluated but rejected due to specific limitations:

- **Edge-betweenness (e.g., Girvan–Newman):** Computationally prohibitive with a complexity of $O(|V||E|^2)$, rendering it impractical for our network ($|V| \approx 50k$, $|E| \approx 20M$).
- **Infomap:** While efficient, it is sensitive to edge weight distributions and tends to generate highly fragmented micro-communities in dense graphs.
- **Label Propagation:** Although extremely fast, its non-deterministic nature leads to unstable partitions, making longitudinal tracking unreliable.

Leiden is a hierarchical algorithm, producing a dendrogram of partitions controlled by the resolution parameter γ.

- Lower values ($\gamma < 1$) favor larger, macro-communities (e.g., broad genre preferences).
- Higher values ($\gamma > 1$) reveal finer micro-communities (e.g., specific niche groups).

In this study, we fixed $\gamma = 1$. This choice provides a balanced granularity, capturing meaningful community structures without merging distinct groups or isolating noise-driven micro-clusters. It also establishes a robust baseline for migration analysis.

To bridge the gap between topology and semantics, we constructed a correlation matrix between the detected user communities and anime genre clusters based on historical rating data. The resulting heatmap, complemented by alluvial flow diagrams, reveals which anime categories dominate within each user group. This approach allows us to uncover community-specific preference profiles and explore the content-driven factors underlying user migration.

6.2 Feature Engineering

To capture the factors driving migration, we constructed a multidimensional feature space describing the user's state at time t. The feature vector $X_u^{(t)}$ comprises four distinct categories:

(1) **Graph Structure (Global & Local):** Measures of centrality and influence, including Degree, Weighted Strength, PageRank, and K-core decomposition.
(2) **Local Cohesion:** This is captured by the weighted clustering coefficient; low values indicate weak neighborhood integration and a potentially higher migration risk.

(3) **Community Embeddedness:** Metrics quantifying the boundary position of a user. Key among these is the *Intra-Community Ratio (ICR)*, defined as the fraction of a user's edges that connect to nodes within the same community.
(4) **Temporal Dynamics:** Delta features representing the year-over-year change in structural metrics (e.g., ΔDegree, ΔICR), capturing the trajectory of a user's engagement.
(5) **Demographics & Activity:** Static attributes (gender, age, location) integrated with dynamic indicators, such as the number of watched titles and rating fluctuations over time.

This results in a comprehensive per-user, per-year feature vector integrating static, structural, and dynamic signals for predictive modeling.

6.3 Model Specification and Performance Assessment

We utilized **CatBoost** [5], a gradient boosting algorithm on decision trees, for the classification task. CatBoost was selected for its native handling of categorical features (e.g., community ID, location) without one-hot encoding, robustness to missing values and different feature scales, and ability to model complex non-linear interactions between structural and behavioral signals. Crucially, it offers native support for class imbalance via internal weighting, allowing the algorithm to effectively handle the rarity of migration events.

Considering the task complexity and strong class imbalance (migration events constitute ≈ 17% of the data), the trained model demonstrates strong predictive performance. It achieves a **ROC AUC of 0.915**, indicating strong discrimination capability.

Table 2: Classification Performance Metrics

Class	Precision	Recall	F1-Score
Non-Migrant (0)	0.93	0.92	0.93
Migrant (1)	0.64	0.68	0.66
Overall Accuracy		0.88	

These results (Table 2) confirm that user migration is not a stochastic process but a predictable phenomenon driven by observable network and behavioral characteristics. This validates the model's utility for early-warning detection, offering threshold adjustments to balance recall and precision.

6.4 Feature Importance and Behavioral Interpretation

To explain the model's decisions, we analyzed SHAP (SHapley Additive exPlanations) summary plot. Figure 4 illustrates the global feature importance and directional impact on the prediction.

The analysis reveals that **Intra-Community Ratio (ICR)** is the dominant predictor. A high ICR acts as a stabilizing force, anchoring the user within their group, whereas a low ICR serves as a strong precursor to migration. Notably, the *temporal decrease* in ICR is the most informative early-warning dynamic signal.

The **current community assignment** itself also has a strong effect, indicating that baseline migration risks differ significantly

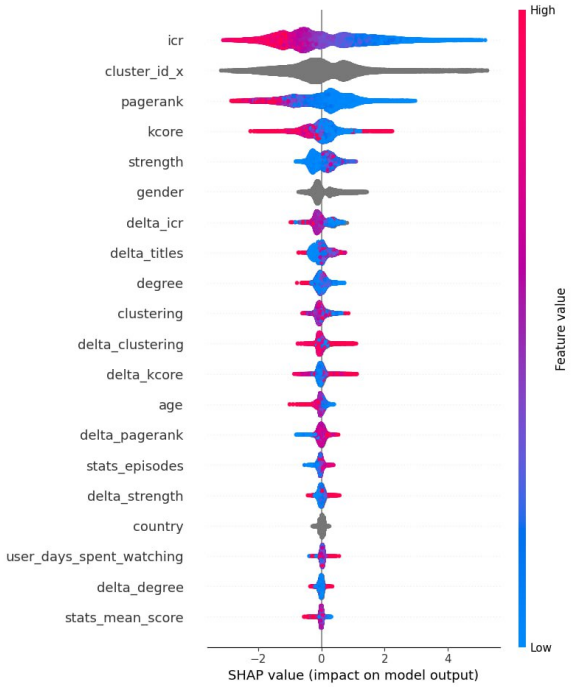

Figure 4: Distribution of SHAP values indicating the contribution of each feature to the model's predictions.

across communities (i.e., some communities are inherently more unstable or "leaky" than others).

PageRank exhibits a non-linear stabilizing effect: highly influential users are less likely to migrate, suggesting that social capital creates "inertia."

Demographically, **older** users show lower migration propensities.

These results emphasize that user migration cannot be captured by singular metrics. Only a comprehensive analysis combining network topology, temporal changes, and behavioral features reveals the true mechanisms driving community shifts.

7 CONCLUSION AND FUTURE WORK

This study presented a comprehensive longitudinal analysis of the MyAnimeList ecosystem (2006–2018), offering a dual perspective on the evolution of digital affinity networks: from the macroscopic topology of the graph to the microscopic dynamics of individual user migration.

Our findings challenge the assumption that expanding social networks necessarily become more cohesive. Instead, we observed a **"Densification-Sparsification Paradox"**: while the user social graph densified into a "Small World," the content interest graph expanded into a sparse "Metropolis" characterized by increasing taste divergence. This confirms that as the anime medium transitioned from a niche subculture to a global entertainment force, the unified community fractured into distinct, self-reinforcing fandoms.

The methodological comparison of community detection algorithms identified the **Leiden algorithm** ($\gamma = 1.0$) as the optimal

tool for longitudinal tracking, striking a balance between structural modularity and semantic interpretability. We demonstrated that "taste" in this domain is a latent variable that correlates with, but is not synonymous with, explicit genre tags.

A critical contribution of this work is the evaluation of user navigation models. The stochastic Random Walker experiment yielded a negative result of significant diagnostic value (similarity $\approx 4\%$), effectively disproving the hypothesis of "Topological Determinism." Users do not passively flow down the gradients of the graph; their trajectories are driven by active, content-aware choices that defy simple Markovian logic.

Consequently, we showed that user migration is a predictable phenomenon only when modeled as a complex interaction between structure and behavior. The CatBoost classifier achieved high accuracy (ROC AUC 0.915), identifying the **Intra-Community Ratio (ICR)** as the dominant predictor of stability. This suggests that the "social gravity" holding a user within a fandom is quantifiable and that migration is preceded by a measurable decay in local ties.

REFERENCES

[1] Azathoth. Anime recommendation database 2020. https://www.kaggle.com/datasets/azathoth42/myanimelist, 2018.

[2] Albert-Laszlo Barabasi and Reka Albert. Emergence of scaling in random networks. *Science*, 286(5439):509–512, 1999.

[3] Everett M. G Borgatti, S. P. Network analysis of 2-mode data. *Social networks*, 19(3):243–264, 1997.

[4] Tamas Nepusz Gabor Csardi. The igraph software package for complex network research. https://igraph.org/, 2005.

[5] Aleksandr Vorobev Anna Veronika Dorogush Andrey Gulin Liudmila Prokhorenkova, Gleb Gusev. Catboost: unbiased boosting with categorical features. In *32nd Conference on Neural Information Processing Systems (NeurIPS 2018)*, 2018.

[6] Renaud Lambiotte Naoki Masuda, Mason A. Porter. Random walks and diffusion on networks. *Physics Reports*, 716(2), 2016.

[7] Iaroslav Sagan. Mars_1.0: Manga and anime research software. https://github.com/BlackSabbitch/MARS_1.0/tree/main/project_cda, 2025. Accessed: 2025-10-27.

[8] Ikvinderpal Singh. Jaccard index versus preferential attachment: A comparative study of similarity based link prediction techniques in complex networks. *Journal of Computer Technology Applications*, 14(2), 2023.

[9] Duncan J. Watts Steven H. Strogatz. Collective dynamics of 'small-world' networks. *Nature*, 393:440–442, 1998.

[10] N. J. van Eck V. A. Traag, L. Waltman. From louvain to leiden: guaranteeing well-connected communities. *Scientific Reports*, 9(5233), 2019.

[11] Hernan Valdivieso. Anime recommendation database 2020. https://www.kaggle.com/datasets/hernan4444/anime-recommendation-database-2020, 2020.